

Argus Incident Response Runbook

This guide describes how the platform operations team handles retrieval service incidents, customer escalations, and post-incident reviews. It is written for engineers who need clear steps during a live issue and enough background to understand why each step exists.

The runbook assumes that telemetry, service ownership, and customer impact records are available in the internal operations console. When those systems are unavailable, the incident commander should keep a manual timeline with exact times, observed symptoms, and decisions that change customer-facing behavior.

Severity Classification

A severity one incident means that production retrieval is unavailable for a broad customer segment, data integrity is at risk, or support cannot provide a workaround. The commander must page the service owner, notify support leadership, and create an incident channel before attempting speculative fixes.

A severity two incident means that retrieval quality or latency is degraded but the service remains usable. The team should identify affected tenants, compare the current error rate with the previous four-hour baseline, and decide whether to pause new deployments until the metric recovers.

A severity three incident covers narrow defects, documentation gaps, or isolated customer reports. These cases still require a ticket, owner, and due date, but they do not require live incident coordination unless the impact grows or duplicates appear from separate customers.

Initial Triage

The first responder checks whether the issue is reproducible with a known query, a known document, and a clean session. The responder should record the query text, source document identifier, retrieval strategy, and the returned chunk identifiers so that later debugging is not based on memory.

If retrieval returns irrelevant chunks, compare the embedding model version, chunking strategy, and index build timestamp. A mismatch between the active index and the latest document version often explains sudden drops in answer quality after a content migration.

If retrieval times out, inspect queue depth, vector store latency, and downstream network errors before changing application code. Most timeout incidents are caused by saturation in a shared dependency rather than by a broken query planner.

Customer Communication

Support updates should be short, factual, and tied to visible progress. The first customer update should confirm that the report was received, describe the observed symptom, and provide the next update time even if the technical cause is still unknown.

Do not promise a resolution time during the first investigation window. A useful message explains what the team is validating, which workaround is available, and what information the customer can provide to accelerate diagnosis.

When a workaround changes retrieval behavior, support must explain the tradeoff. For example, increasing chunk size can improve recall for long policy documents but may increase response latency and the amount of context sent to the answer generator.

Rollback and Mitigation

Rollback is preferred when a recent deployment correlates with the first bad metric and the previous version is known to be healthy. The rollback owner should verify that configuration, database migrations, and index schema changes remain compatible with the older service version.

Mitigation is preferred when the defect is caused by input data, external dependency behavior, or a workload spike. Common mitigations include reducing top-k retrieval depth, disabling optional reranking, rebuilding only affected tenant indexes, or temporarily routing traffic to a smaller document collection.

Every mitigation must include an exit condition. The team should define the metric threshold, observation window, and owner responsible for removing temporary limits after the service returns to normal.

Post-Incident Review

The review document should include customer impact, detection path, timeline, contributing factors, and follow-up actions. Avoid blame-oriented language; focus on the system conditions that made the failure possible or made detection slower than expected.

Action items should be specific enough that a reviewer can verify completion without private context. Good actions name the service, metric, alert, dashboard, test, or runbook section that will change, and they include an owner and due date.

If evaluation samples failed to catch the incident, add representative queries and evidence spans to the benchmark suite. The benchmark should test the retrieval behavior that customers actually needed, not only the chunking strategy that looked cheapest during development.

Benchmark Notes

For retrieval benchmarks, compare recall, reciprocal rank, normalized discounted cumulative gain, estimated embedding token cost, and candidate latency together. A strategy that is cheap but repeatedly misses evidence should not be recommended for production use.

Paragraph chunking usually performs well when documents have clean headings and blank lines. Recursive character chunking is helpful when sections contain long paragraphs, tables converted to text, or mixed formatting where simple paragraph boundaries become unreliable.

Manual inspection remains necessary because synthetic queries can be too easy, too repetitive, or tied to a single section of the document. At least three evidence samples should be checked before trusting the recommendation summary.